



Processes

How can many people share one computer efficiently?




Multiprocessing



- **Sharing through hardware abstraction**
 - Software(OS) managed hardware(CPU, memory, devices)
 - Providing **consistent abstraction** view to every user program
- Computer runs many processes simultaneously
 - Applications for one or more users
 - Web browsers, editors, ...
 - Background tasks
 - Monitoring network & I/O devices


Multiprocessing Example



- Running program “htop” on department machine
 - System has 69 “tasks”, 1 of which is running
 - Identified by Process ID (PID), user account, command name

Processes


- Definition: A *process* is an instance of a running program.
 - Not the same as “program” or “processor”
- Process provides each program with two key abstractions:
 - ***Private address space***
 - Each program seems to have exclusive use of main memory.
 - Provided by kernel mechanism called *virtual memory*
 - ***Logical control flow***
 - Each program seems to have exclusive use of the CPU
 - Provided by kernel mechanism called *context switching*



Control Flow

- Processors do only one thing:
 - From startup to shutdown, each CPU core simply reads and executes a sequence of machine instructions, one at a time *
 - This sequence is the CPU's *control flow* (or *flow of control*)


Physical control flow




* many modern CPUs execute several instructions at once and/or out of program order, but this is invisible to the programmer

Context Switching

- Processes are managed by a shared chunk of memory-resident OS code called the *kernel*
 - Important: the kernel is not a separate process, but rather runs as part of some existing process.
- Control flow passes from one process to another via a *context switch*




Context Switching (Uniprocessor)




- Single processor executes multiple processes concurrently
 - Process executions **interleaved** (multitasking)
 - Address spaces managed by virtual memory system
 - Register values for nonexecuting processes saved in memory

Context Switching (Uniprocessor)




- Save current registers in memory

Context Switching (Uniprocessor)




- Schedule next process for execution

Context Switching (Uniprocessor)



- Load saved registers and switch address space (context switch)


Context Switching (Multicore)



- Multicore processors
 - Multiple CPUs on single chip
 - Share main memory (and some caches)
 - Each can execute a separate process
 - Scheduling of processes onto cores done by kernel


Uniprocessor Logical control flows

- Only one process runs at a time
- A and B execution is *interleaved*, not truly concurrent
- Similarly for A and C
- Still possible for A and B / A and C to interfere with each other



Concurrent flows


- Two processes *run concurrently* (are concurrent) if their execution overlaps in time
- Otherwise, they are *sequential*
- Appears as if concurrent processes run in parallel with each other
 - This means they can interfere with each other



How does the kernel take control?

- The CPU executes instructions in sequence
- We don't write "now run kernel code" in our programs...
 - Timer interrupts: generate by hardware timer
 - *Or do we??*

Physical control flow



System Calls

- Whenever a program wants to cause an effect outside its own process, it must ask the kernel for help

- Examples:

- Read/write files
- Get current time
- Allocate RAM (sbrk)
- Create new processes

```
1 // fopen.c
2 FILE *fopen(const char *fname, const char *mode)
3 {
4     int flags = mode2flags(mode);
5     if (!flags) return NULL;
6     int fd = open(fname, flags, DEFLPERMS);
7     if (fd == -1) return NULL;
8     return fdopen(fd, mode);
9 }
10
11 // open.S
12     .global open
13 open:
14     mov $SYS_open, %eax
15     syscall
16     cmp $SYS_error_thresh, %rax
17     ja __syscall_error
18     ret
```

System Call Error Handling

- Almost all system-level operations can fail
 - Only exception is the handful of functions that return `void`
 - You must explicitly check for failure
- On error, most system-level functions return `-1` and set global variable `errno` to indicate cause.
- Example:

```
1 pid_t pid = fork();
2 if (pid == -1) {
3     fprintf(stderr, "fork error: %s\n", strerror(errno));
4     exit(1);
5 }
```

Obtaining Process IDs

- Every process has a process ID, a code number that identifies it
- PID does not change during the process lifetime
- Every process has a parent process
 - The process that started it
 - Except the [init](#) process(pid 1)
- `pid_t getpid(void)`
 - Returns PID of current process
- `pid_t getppid(void)`
 - Returns PID of parent process

Process States

At any time, each process is either:

- **Running**
 - Process is either executing instructions, or it *could be* executing instructions if there were enough CPU cores.
- **Blocked / Sleeping**
 - Process cannot execute any more instructions until some external event happens (usually I/O).
- **Stopped**
 - Process has been prevented from executing by user action (control-Z).
- **Terminated / Zombie**
 - Process is finished. Parent process has not yet been notified.

Terminating Processes

- Process becomes terminated for one of three reasons:
 - Receiving a signal whose default action is to terminate
 - Returning from the `main` routine
 - Calling the `exit` function
- `void exit(int status)`
 - Terminates with an *exit status* of `status`
 - Convention: normal return status is 0, nonzero on error
 - Another way to explicitly set the exit status is to return an integer value from the main routine
- `exit` is called **once** but **never** returns.

Creating Processes


- *Parent process* creates a new running *child process* by calling `fork`
- `int fork(void)`
 - Returns 0 to the child process, child's **PID** to parent process
 - The way to distinguish the parent process and the child process
 - Child is *almost* identical to parent:
 - Child get an **identical** (but separate) copy of the parent's virtual address space.
 - Child gets **identical** copies of the parent's open file descriptors
 - Child has a **different** PID than the parent
- `fork` is interesting (and often confusing) because it is called **once** but returns **twice**

Differences between parent and child


- The return values from fork are different.
- The process IDs are different (`getpid()` to obtain process ID).
- The two processes have different parent process IDs (`getppid()` to obtain parent process ID).
- File locks set by the parent are not inherited by the child.
- Pending alarms are cleared for the child.
- The set of pending signals for the child is set to the empty set.

Parent process and child process almost identical

- The two processes have
 - the **same program text**.
 - but **separate copies of the data, stack, and heap segments**.
- The child's data, stack, and heap segments are **initially** exact duplicates of the corresponding parts the parent's memory.
- After the `fork()`, each process can modify the variables in its data, stack, and heap segments without affecting the other process.



Conceptual View of fork



- Make complete copy of execution state
 - Designate one as parent and one as child
 - Resume execution of parent or child

Memory semantics of fork()

- Problems of direct copying parent's (virtual) memory pages when forking.
 - **Slow process creation** - the child process won't be ready until all the memory copying is done.
 - **Waste of resources and time**
 - Many resources could be **shared** between parent and child.
 - fork() is usually immediately followed by an exec function to run a different program
 - Causes a *replacement* of the child process text segment, and *reinitialization* of the data, stack and heap segments, making the direct copying when forking a **waste**

Memory semantics of fork()


Most modern Unix implementations use two techniques to avoid the wasteful copying:

- For the **text segment** of each process, the kernel marks it as **read-only**.
 - A process cannot modify its own code.
 - When fork() creates text segment of the child, it builds the page table entries for the text segment that have **the same virtual/physical mappings** as the parent.
- Applying the **copy-on-write** technique for data, stack, and heap segments

Copy-on-write for fork()

- After fork(), the child and parent **share the same memory mappings**
- Memory pages (data, stack, heap) are **marked as read-only** for both processes
- When either process tries to **modify a page**, the kernel
 - Detects the write attempt
 - Creates a **duplicate copy** of the page for the modifying process
 - Updates the process's page table with the new copy
- **Both processes now have separate writable copies** of the modified pages
- **This avoids unnecessary duplication** and improves performance

Before modification



fork Example

```
1 int main(int argc, char** argv)
2 {
3     pid_t pid;
4     int x = 1;
5
6     pid = fork();
7     if (pid == 0) { /* Child */
8         printf("child : x=%d\n", ++x);
9         return 0;
10    }
11
12    /* Parent */
13    printf("parent: x=%d\n", --x);
14    return 0;
15 }
```

- Call once, return twice

- Concurrent execution

- Can't predict **execution order** of parent and child

```
linux> ./fork
parent: x=0
child : x=2
```

```
linux> ./fork
child : x=2
parent: x=0
```

```
linux> ./fork
parent: x=0
child : x=2
```


```
linux> ./fork
parent: x=0
child : x=2
```

Modeling fork with Process Graphs

- Represents partial ordering of statements in a concurrent program
- Vertices = Execution of statements
- Edges ($a \rightarrow b$) = a happens before b
- Labels on edges may show variable values.
- `printf` vertices indicate output
- Starts with a vertex having no incoming edges
- Topological sort gives a valid total execution order


Process Graph Example

```
1 int main(int argc, char** argv)
2 {
3     pid_t pid;
4     int x = 1;
5
6     pid = fork();
7     if (pid == 0) { /* Child */
8         printf("child : x=%d\n", ++x);
9         return 0;
10    }
11
12    /* Parent */
13    printf("parent: x=%d\n", --x);
14    return 0;
15 }
```




Interpreting Process Graphs


- Original graph:




- Relabeled graph:



Feasible total ordering:




Feasible or Infeasible?



Infeasible: not a topological sort

fork Example: Two consecutive forks

```
1 void fork2()
2 {
3     printf("L0\n");
4     fork();
5     printf("L1\n");
6     fork();
7     printf("Bye\n");
8 }
```



Feasible

L0

L1

Bye

Bye

L1

Bye

Bye

Infeasible

L0

Bye

L1

Bye


L1

Bye

Bye

fork Example: Nested forks in parent

```
1 void fork4()
2 {
3     printf("L0\n");
4     if (fork() != 0) {
5         printf("L1\n");
6         if (fork() != 0) {
7             printf("L2\n");
8         }
9     }
10    printf("Bye\n");
11 }
```



Feasible or Infeasible?

L0

Bye

L1

Bye

Bye

L2

Infeasible

Feasible or Infeasible?

L0

L1

Bye

Bye


L2

Bye

Feasible

fork Example: Nested forks in children

```
1 void fork5()
2 {
3     printf("L0\n");
4     if (fork() == 0) {
5         printf("L1\n");
6         if (fork() == 0) {
7             printf("L2\n");
8         }
9     }
10    printf("Bye\n");
11 }
```



Feasible or Infeasible?

L0

Bye

L1

Bye

Bye

L2

Infeasible

Feasible or Infeasible?

L0

Bye

L1

L2

Bye

Bye

Feasible

Quiz

After a program executes the following series of fork() calls, how many new processes will result (assuming that none of the calls fails)?

```
fork();  
fork();  
fork();
```

The vfork() function

```
#include <unistd.h>

pid_t vfork(void);
```

In parent: returns process ID of child on success, or -1 on error;
in successfully created child: always returns 0

- vfork() function is used in the same way as fork()
- vfork() is generally designed to be used in programs where the child performs an immediate exec call

The vfork() function

- Difference between vfork() from fork(), vfork () has the following features
 - **No duplication** of parent process address space.
 - The child **shares** the parent's memory until it either performs a successful execve() or calls _exit() to terminate.
 - Execution of the parent process is **suspended** until the child has performed an exec() or _exit()
 - The child process is **always** scheduled to run **before** the parent process with vfork() (but not true for fork())
 - vfork() is more **efficient** since it does not make a copy of parent's address space, however, be careful to avoid modification of parent's address space

The vfork() function

- The `vfork()` function originated with early BSD Unix, whose implementation of `fork()` is performing direct copy after forking (i.e., **without COW**)
- BSD introduced `vfork()` to solve the inefficiency problem of `fork()`
- Now modern Unix implementations have COW for `fork()`, thus largely eliminating the need for `vfork()`
- Linux (and many other UNIX implementations) provide a `vfork()` system call with BSD semantics for programs that **require the fastest possible fork**
 - Except where speed is absolutely critical, new programs should avoid the use of `vfork()` in favor of `fork()`

Reaping Child Processes

- Idea

- When process terminates, it **still consumes** system resources, kept around in a terminated state
 - Examples: Exit status, various OS tables
- Called a “**zombie**” process
 - Terminated but not yet reaped

- Reaping

- Performed by parent on terminated child (using `wait` or `waitpid`)
- Parent is given exit status information
- Kernel **then** deletes zombie child process

- What if parent doesn't reap?

- If any parent terminates without reaping a child, then the **orphaned** child should be reaped by **init** process (`pid == 1`)
 - Unless it was `init` that terminated! Then need to reboot...
- Only need explicit reaping in long-running processes
 - e.g., shells and servers

Zombie Example

```
linux> ./forks 7 &
[1] 6639
Running Parent, PID = 6639
Terminating Child, PID = 6640
```

```
linux> ps
  PID TTY      TIME CMD
 6585 ttyp9    00:00:00 tcsh
 6639 ttyp9    00:00:03 forks
 6640 ttyp9    00:00:00 forks <defunct>
 6641 ttyp9    00:00:00 ps
linux> kill 6639
[1]  Terminated
```

```
linux> ps
  PID TTY      TIME CMD
 6585 ttyp9    00:00:00 tcsh
 6642 ttyp9    00:00:00 ps
```

```
1 void fork7() {
2     if (fork() == 0) {
3         /* Child */
4         printf("Terminating Child, PID = %d\n", getpid());
5         exit(0);
6     } else {
7         printf("Running Parent, PID = %d\n", getpid());
8         while (1)
9             ; /* Infinite loop */
10    }
11 }
```

- ps shows child process as “defunct” (i.e., a zombie)

- Killing parent allows child to be reaped by init

Non-terminating Child Example

```
linux> ./forks 8
Terminating Parent, PID = 6675
Running Child, PID = 6676
```

```
linux> ps
  PID TTY      TIME CMD
 6585 ttyp9    00:00:00 tcsh
 6676 ttyp9    00:00:06 forks
 6677 ttyp9    00:00:00 ps
```

```
linux> kill 6676 ←
linux> ps
  PID TTY      TIME CMD
 6585 ttyp9    00:00:00 tcsh
 6678 ttyp9    00:00:00 ps
```

```
1 void fork8()
2 {
3     if (fork() == 0) {
4         /* Child */
5         printf("Running Child, PID = %d\n",
6                getpid());
7         while (1)
8             ; /* Infinite loop */
9     } else {
10        printf("Terminating Parent, PID = %d\n",
11               getpid());
12        exit(0);
13    }
14 }
```


- Child process still active even though parent has terminated
- Must kill child explicitly, or else will keep running indefinitely

wait: Synchronizing with Children

- Parent **reaps** a child process with one of these system calls:
- `pid_t wait(int *status)`
 - Suspends current process until one of its children terminates
 - Returns PID of child, records exit status in **status**
- `pid_t waitpid(pid_t pid, int *status,
int options)`
 - More flexible version of **wait**
 - Can wait for a specific child or group of children
 - Can be told to return immediately if there are no children to reap

wait: Synchronizing with Children

```
1 void fork9() {  
2     int child_status;  
3  
4     if (fork() == 0) {  
5         printf("HC: hello from child\n");  
6         exit(0);  
7     } else {  
8         printf("HP: hello from parent\n");  
9         wait(&child_status);  
10        printf("CT: child has terminated\n");  
11    }  
12    printf("Bye\n");  
13 }
```



Feasible output(s):

HC HP

HP HC

CT CT

Bye Bye

Infeasible output:

HP

CT

Bye

HC

wait: Status codes

- Return value of `wait` is the pid of the child process that terminated
- If `status != NULL`, then the integer it points to will be set to a value that indicates the exit status
 - More information than the value passed to `exit`
 - Must be decoded, using macros defined in `sys/wait.h`
 - `WIFEXITED`, `WEXITSTATUS`, `WIFSIGNALED`, `WTERMSIG`, `WIFSTOPPED`, `WSTOPSIG`, `WIFCONTINUED`

Another wait Example

- If multiple children completed, will take in arbitrary order
- Can use macros WIFEXITED and WEXITSTATUS to get information about exit status

```
1 void fork10() {
2     pid_t pid[N];
3     int i, child_status;
4
5     for (i = 0; i < N; i++)
6         if ((pid[i] = fork()) == 0) {
7             exit(100+i); /* Child */
8         }
9     for (i = 0; i < N; i++) { /* Parent */
10        pid_t wpid = wait(&child_status);
11        if (WIFEXITED(child_status))
12            printf("Child %d terminated with exit status %d\n",
13                  wpid, WEXITSTATUS(child_status));
14        else
15            printf("Child %d terminate abnormally\n", wpid);
16    }
17 }
```

waitpid: Waiting for a Specific Process

- pid_t waitpid(pid_t pid, int *status, int options)
- **Suspends** current process until specific process terminates

```
1 void fork11() {  
2     pid_t pid[N];  
3     int i;  
4     int child_status;  
5  
6     for (i = 0; i < N; i++)  
7         if ((pid[i] = fork()) == 0)  
8             exit(100+i); /* Child */  
9     for (i = N-1; i >= 0; i--) {  
10         pid_t wpid = waitpid(pid[i], &child_status, 0);  
11         if (WIFEXITED(child_status))  
12             printf("Child %d terminated with exit status %d\n",  
13                     wpid, WEXITSTATUS(child_status));  
14         else  
15             printf("Child %d terminate abnormally\n", wpid);  
16     }  
17 }
```

sleep: putting processes to sleep

- `unsigned int sleep(unsigned int secs);`
 - Put process into sleep state
 - returns 0 if the requested amount of time has elapsed
 - Or number of seconds still left to sleep


- `int pause(void);`
 - Puts the calling process to sleep until a signal is received by the process

`execve`: Loading and Running Programs

- `int execve(char *filename, char *argv[], char *envp[])`
- Loads and runs in the current process:
 - Executable file `filename`
 - (e.g., `#!/bin/bash`) Can be object file or script file beginning with `#!interpreter`
 - ...with argument list `argv`
 - By convention `argv[0]==filename`
 - ...and environment variable list `envp`
 - “name=value” strings
- **Overwrites** code, data, heap, and stack segments
 - Retains PID, open files and signal context
- Called **once** and **never** returns
 - ...except if there is an error


execve Example

- Execute "/bin/ls -lt /usr/include" in child process using current environment:



```
1 if ((pid = fork()) == 0) { /* Child runs program */
2     if (execve(myargv[0], myargv, environ) < 0) {
3         printf("%s: %s\n", myargv[0], strerror(errno));
4         exit(1);
5     }
6 }
```

execve and process memory layout



- To load and run a new program `a.out` in the current process using `execve`
- Free `vm_area_struct`'s and page tables for old areas
- Create `vm_area_struct`'s and page tables for new areas
 - Programs and initialized data backed by object files.
- Set PC to entry point in `.text`
 - Linux will fault in code and data pages as needed

fork and execve to run programs

- Programs like shells and Web servers make heavy use of fork and execve
 - Traditional web server for handling connections, or to show dynamic contents
 - Use shell to run commands or program

Simple Shell Implementation

- Basic loop
 - Read line from command line
 - Execute the requested operation
 - Built-in command (only one implemented is `quit`)
 - Load and execute program from file

```
1 int main(int argc, char** argv)
2 {
3     char cmdline[MAXLINE]; /* command line */
4
5     while (1) {
6         /* read */
7         printf("> ");
8         fgets(cmdline, MAXLINE, stdin);
9         if (feof(stdin))
10             exit(0);
11
12         /* evaluate */
13         eval(cmdline);
14     }
15 }
```

Simple Shell eval Function

```
void eval(char * cmdline)
{
    char * argv[MAXARGS]; /* Argument list execve() */
    char buf[MAXLINE]; /* Holds modified command line */
    int bg; /* Should the job run in bg or fg? */
    pid_t pid; /* Process id */

    strcpy(buf, cmdline);
    bg = parseline(buf, argv);           parseline will parse 'buf' into 'argv' and return whether or
    if (argv[0] == NULL)                  not input line ended in '&'
        return; /* Ignore empty lines */
}
```

Simple Shell eval Function

```
void eval(char * cmdline)
{
    char * argv[MAXARGS]; /* Argument list execve() */
    char buf[MAXLINE]; /* Holds modified command line */
    int bg; /* Should the job run in bg or fg? */
    pid_t pid; /* Process id */

    strcpy(buf, cmdline);
    bg = parseline(buf, argv);
    if (argv[0] == NULL)
        return; /* Ignore empty lines */
```

Ignore empty lines.

Simple Shell eval Function

```
void eval(char * cmdline)
{
    char * argv[MAXARGS]; /* Argument list execve() */
    char buf[MAXLINE]; /* Holds modified command line */
    int bg; /* Should the job run in bg or fg? */
    pid_t pid; /* Process id */

    strcpy(buf, cmdline);
    bg = parseline(buf, argv);
    if (argv[0] == NULL)
        return; /* Ignore empty lines */

    if (!builtin_command(argv)) {
```

If it is a 'built in' command, then handle it here in this program. Otherwise fork/exec the program specified in argv[0]

Simple Shell eval Function

```
void eval(char * cmdline)
{
    char * argv[MAXARGS]; /* Argument list execve() */
    char buf[MAXLINE]; /* Holds modified command line */
    int bg; /* Should the job run in bg or fg? */
    pid_t pid; /* Process id */

    strcpy(buf, cmdline);
    bg = parseline(buf, argv);
    if (argv[0] == NULL)
        return; /* Ignore empty lines */

    if (!builtin_command(argv)) {
        if ((pid = fork()) == 0) { /* Child runs user job */
            Create child
        }
    }
}
```

Create child

Simple Shell eval Function

```
void eval(char * cmdline)
{
    char * argv[MAXARGS]; /* Argument list execve() */
    char buf[MAXLINE]; /* Holds modified command line */
    int bg; /* Should the job run in bg or fg? */
    pid_t pid; /* Process id */

    strcpy(buf, cmdline);
    bg = parseline(buf, argv);
    if (argv[0] == NULL)
        return; /* Ignore empty lines */

    if (!builtin_command(argv)) {
        if ((pid = fork()) == 0) { /* Child runs user job */
            execve(argv[0], argv, environ);
            // If we get here, execve failed.
            printf("%s: %s\n", argv[0], strerror(errno));
            exit(127);
        }
    }
}
```

Start argv[0].
execve only returns on error.
127 used for command not found

Simple Shell eval Function

```
void eval(char * cmdline)
{
    char * argv[MAXARGS]; /* Argument list execve() */
    char buf[MAXLINE]; /* Holds modified command line */
    int bg; /* Should the job run in bg or fg? */
    pid_t pid; /* Process id */

    strcpy(buf, cmdline);
    bg = parseline(buf, argv);
    if (argv[0] == NULL)
        return; /* Ignore empty lines */

    if (!builtin_command(argv)) {
        if ((pid = fork()) == 0) { /* Child runs user job */
            execve(argv[0], argv, environ);
            // If we get here, execve failed.
            printf("%s: %s\n", argv[0], strerror(errno));
            exit(127);
        }

        /* Parent waits for foreground job to terminate */
        if (!bg) {
            int status;
            if (waitpid(pid, &status, 0) < 0)
                unix_error("waitfg: waitpid error");
        }
    }
}
```

If running child in foreground, wait until it is done.

Simple Shell eval Function

```
void eval(char * cmdline)
{
    char * argv[MAXARGS]; /* Argument list execve() */
    char buf[MAXLINE]; /* Holds modified command line */
    int bg; /* Should the job run in bg or fg? */
    pid_t pid; /* Process id */

    strcpy(buf, cmdline);
    bg = parseline(buf, argv);
    if (argv[0] == NULL)
        return; /* Ignore empty lines */

    if (!builtin_command(argv)) {
        if ((pid = fork()) == 0) { /* Child runs user job */
            execve(argv[0], argv, environ);
            // If we get here, execve failed.
            printf("%s: %s\n", argv[0], strerror(errno));
            exit(127);
        }

        /* Parent waits for foreground job to terminate */
        if (!bg) {
            int status;
            if (waitpid(pid, &status, 0) < 0)
                unix_error("waitfg: waitpid error");
        }
        else
            printf("%d %s\n", pid, cmdline);
    }
    return;
}
```

If running child in background,
print pid and continue doing
other stuff.

Problem with Simple Shell Example

- Shell designed to run **indefinitely**
 - Should not accumulate unneeded resources
 - Memory
 - Child processes
 - File descriptors
- Our example shell correctly waits for and reaps foreground jobs
- But what about **background jobs**?
 - Will become zombies when they terminate
 - Will never be reaped because shell (typically) will not terminate
 - Could run the entire computer out of memory
 - Moreover, run out of PIDs

Summary

- Processes
 - At any given time, system has multiple active processes
 - Only one can execute at a time on any single core
 - Each process *appears* to have total control of processor + private memory space

Summary (cont.)

- Spawning processes
 - Call `fork`
 - One call, two returns
- Process completion
 - Call `exit`
 - One call, no return
- Put process into sleep
 - Call `sleep`, `pause`
- Reaping and waiting for processes
 - Call `wait` or `waitpid`
- Loading and running programs
 - Call `execve` (or variant)
 - One call, (normally) no return